

A data-driven distributionally robust bound on the expected optimal value of uncertain mixed 0-1 linear programming

Guanglin Xu¹ · Samuel Burer¹

Received: 13 September 2017 / Accepted: 6 January 2018 / Published online: 17 January 2018
© Springer-Verlag GmbH Germany, part of Springer Nature 2018

Abstract This paper studies the expected optimal value of a mixed 0-1 programming problem with uncertain objective coefficients following a joint distribution. We assume that the true distribution is not known exactly, but a set of independent samples can be observed. Using the Wasserstein metric, we construct an ambiguity set centered at the empirical distribution from the observed samples and containing the true distribution with a high statistical guarantee. The problem of interest is to investigate the bound on the expected optimal value over the Wasserstein ambiguity set. Under standard assumptions, we reformulate the problem into a copositive program, which naturally leads to a tractable semidefinite-based approximation. We compare our approach with a moment-based approach from the literature on three applications. Numerical results illustrate the effectiveness of our approach.

Keywords Distributionally robust optimization · Wasserstein metric · Copositive programming · Semidefinite programming

1 Introduction

We consider the following uncertain mixed 0-1 linear programming problem:

$$v(\xi) := \max \left\{ (F\xi)^T x : \begin{array}{l} Ax = b, x \geq 0 \\ x_j \in \{0, 1\} \forall j \in \mathcal{B} \end{array} \right\} \quad (1)$$

✉ Guanglin Xu
guanglin-xu@uiowa.edu

Samuel Burer
samuel-burer@uiowa.edu

¹ Department of Management Sciences, University of Iowa, Iowa City, IA 52242-1994, USA

where $A \in \mathbb{R}^{m \times n}$, $F \in \mathbb{R}^{n \times k}$, and $b \in \mathbb{R}^m$ are the problem data, $x \in \mathbb{R}_+^n$ is the vector of decision variables, $\mathcal{B} \subseteq \{1, \dots, n\}$ is an index set of binary variables, and the objective coefficients are linear in the random vector $\xi \in \mathbb{R}^k$ via F . Problem (1) entails two extreme classes of programs: if $\mathcal{B} = \emptyset$, then (1) represents the regular linear program with uncertain objective coefficients; if $\mathcal{B} = \{1, \dots, n\}$, then (1) represents the regular binary program with uncertain coefficients. In general, problem (1) is NP-hard (Wolsey 1998).

The optimal value $v(\xi)$ is a random variable as ξ is a random vector. We assume that ξ follows a multivariate distribution \mathbb{P} supported on a nonempty set $\Xi \subseteq \mathbb{R}^k$, which is, in particular, defined as a slice of a closed, convex, full-dimensional cone $\widehat{\Xi} \subseteq \mathbb{R}_+ \times \mathbb{R}^{k-1}$:

$$\Xi := \left\{ \xi \in \widehat{\Xi} : e_1^T \xi = \xi_1 = 1 \right\},$$

where e_1 is the first standard basis vector in \mathbb{R}^k . In words, $\widehat{\Xi}$ is the homogenization of Ξ . We choose this homogenized version for notational convenience. Note that it, in fact, enables us to model affine effects of the uncertain parameters in (1).

The expected optimal value of (1), denoted by $v_{\mathbb{P}}$, is defined as

$$v_{\mathbb{P}} := \mathbb{E}_{\mathbb{P}}[v(\xi)] = \int_{\Xi} v(\xi) d\mathbb{P}(\xi).$$

The problem of computing $v_{\mathbb{P}}$ has been extensively studied in the literature. Hagstrom (1988) showed that computing $v_{\mathbb{P}}$ for the longest path problem over a directed acyclic graph is $\#\mathcal{P}$ -complete even if the arc lengths are each independently distributed and restricted to taking two possible values. Aldous (2001) studied a linear assignment problem with random cost coefficients following either an independent uniform distribution on $[0, 1]$ or an exponential distribution with parameter 1 and proved that the asymptotic value of $v_{\mathbb{P}}$ approaches $\frac{\pi^2}{6}$ as the number of assignments goes to infinity. For additional studies, see (Bowman 1995; Dodin 1985; Möhring 2001).

In practice, it is difficult or impossible to know \mathbb{P} completely, and computing $v_{\mathbb{P}}$ is thus not well defined in this situation. An alternative is to construct an ambiguity set, denoted by \mathcal{D} , that contains a family of distributions supported on Ξ and consistent with any known properties of \mathbb{P} . Ideally, the ambiguity set will possess some statistical guarantee, e.g., the probability that $\mathbb{P} \in \mathcal{D}$ will be at least $1 - \beta$, where β is the significance level. In analogy with $v_{\mathbb{P}}$, we define $v_{\mathbb{Q}}$ for any $\mathbb{Q} \in \mathcal{D}$. Then, we are interested in computing the maximum expected optimal value $v_{\mathbb{Q}}$ over the ambiguity set \mathcal{D} :

$$v_{\mathcal{D}}^+ := \sup_{\mathbb{Q} \in \mathcal{D}} v_{\mathbb{Q}}. \quad (2)$$

Note that, when the probability of $\mathbb{P} \in \mathcal{D}$ is at least $1 - \beta$, the probability of $v_{\mathbb{P}} \leq v_{\mathcal{D}}^+$ is at least $1 - \beta$.

There are three main issues (somehow conflicting) regarding the computation of $v_{\mathcal{D}}^+$. First, one would like an ambiguity set \mathcal{D} with a high statistical guarantee to contain the true distribution \mathbb{P} . In this way, the computed $v_{\mathcal{D}}^+$ will be an upper bound on $v_{\mathbb{P}}$

with a high confidence level. (We will introduce several approaches in the following paragraph.) Second, one would like $v_{\mathcal{D}}^+$ to be tight in the sense that it is as close to $v_{\mathbb{P}}$ as possible. Generally, if \mathcal{D} enforces more information about \mathbb{P} , then $v_{\mathcal{D}}^+$ will be closer to $v_{\mathbb{P}}$. Finally, the third concern is the complexity of the resulting optimization problem, i.e., whether the problem can be solved in polynomial time.

Bertsimas et al. (2004, 2006) constructed moment ambiguity sets using the first two marginal moments of each ξ_i . Denote the first and second of each uncertain parameter by μ_i and σ_i respectively. They computed $v_{\mathcal{D}}^+$ over all joint distributions sharing the same first two marginal moments and proved polynomial-time computability if the corresponding deterministic problem is solvable in polynomial time. However, the computed bound may not be tight with respect to $v_{\mathbb{P}}$ since the marginal-moment model does not capture the dependence of the random variables. In a closely related direction, Natarajan et al. (2009) proposed an ambiguity set that was constructed from the known marginal distributions of each random variable ξ_i , and they computed $v_{\mathcal{D}}^+$ by solving a concave maximization problem. As an extension to the marginal moment-based approach, Natarajan et al. (2011) proposed a cross-moment model that was based on an ambiguity set constructed using both marginal and cross moments. Compared to the marginal-moment approach, the cross-moment approach has tighter upper bounds as the model captures the dependence of the random variables. However, computing the bound requires solving a completely positive program, which itself can only be approximated in general. Thus, the authors proposed semidefinite programming (SDP) relaxations to approximate $v_{\mathcal{D}}^+$.

Moment-based ambiguity sets are also used prominently in a parallel vein of research, called distributionally robust optimization (DRO); see (Bertsimas and Popescu 2005; Calafiore and Ghaoui 2006; Chen et al. 2016; Delage and Ye 2010; Dupačová 1987; Erdoğan and Iyengar 2006; Ghaoui et al. 2003; Hanasusanto et al. 2015; Natarajan and Teo 2017; Vandenberghe et al. 2007; Zymler et al. 2013a,b). The popularity of the moment-based approach is mainly due to the fact that it often leads to tractable optimization problems and relatively simple models. Its weakness, however, is that moment-based sets are not guaranteed to converge to the true distribution \mathbb{P} when the sample size increases to infinity, even though the estimations of the first and second moments are themselves guaranteed to converge.

As an attractive alternative to moment-based ambiguity sets, distance-based ambiguity sets have been proposed in recent years. This approach defines \mathcal{D} as a ball in the space of probability distributions equipped with a distance measure, and the center of the ball is typically the empirical distribution derived from a series of independent realizations of the random vector ξ . The key ingredient of this approach is the distance function. Classical distance functions include the Kullback–Leibler divergence (Hu and Hong 2013; Jiang and Guan 2016), the ϕ -divergence (Ben-Tal et al. 2013; Duchi et al. 2016; Klabjan et al. 2013), the Prohorov metric (Erdoğan and Iyengar 2006), empirical Burg-entropy divergence balls (Lam 2016), and the Wasserstein metric (Pflug and Wozabal 2007; Wiesemann et al. 2014).

In this paper, we apply the Wasserstein metric to construct a data-driven ambiguity set \mathcal{D} centered at the empirical distribution $\hat{\mathbb{P}}_N$ derived from N independent observations of ξ . This approach has several benefits. The conservativeness of the ambiguity

set can be controlled by tuning a single parameter, the radius of the Wasserstein ball; we will discuss this parameter in detail in Sect. 2. Also, under mild conditions on \mathbb{P} , the Wasserstein ambiguity provides a natural confidence set for \mathbb{P} . Specifically, the Wasserstein ball around the empirical distribution on N independent identical samples contains \mathbb{P} with confidence $1 - \beta$ if its radius exceeds an explicit threshold $\epsilon_N(\beta)$ that can be computed via a closed form equation (Eichfelder and Jahn 2008; Esfahani and Kuhn 2015). We then formulate v_D^+ in (2) over the constructed Wasserstein ambiguity set. That is, we model the maximum value of v_Q over the ambiguity set \mathcal{D} constructed by the Wasserstein metric. In Sect. 3, we reformulate problem (2) into a copositive problem under some standard assumptions. As the copositive reformulation is computationally intractable, we apply a standard approach based on semidefinite programming techniques to approximate v_D^+ from above. In Sect. 4, we numerically verify our approach on three applications from the literature. In particular, we compare our approach with the moment-based approach proposed in Natarajan et al. (2011). We have several important observations from the experimental results. First, we find that the gaps between the bound from our semidefinite programs and the true expected optimal value becomes narrower as the sample size increases. However, the moment-based bound remains the same regardless of the increase in the sample size. Second, we observe that our bound converges to the true expected optimal value on the first two applications where the underlying deterministic problems are linear programs. Although our bound on the third application is not able to converge to the true expected optimal value, it is tighter than the moment-based bound after the sample size increases to a certain level. We conclude our research and discuss some future directions in Sect. 5.

We point out some similarities of our paper to a recent technical report by Hanasusanto and Kuhn (2016). In their report, they proposed a Wasserstein-metric ambiguity set for a two-stage DRO problem. In particular, they applied copositive programming techniques to reformulate the second-stage worst-case value function, which is essentially a max-min optimization problem, while we use copositive techniques to reformulate a max-max optimization problem; see (4). Furthermore, they directly used a hierarchy schema to approximate the copositive cones, while we derive natural SDP approximations based on the copositive reformulation. Note that their hierarchy of approximations lead to SDP approximations as well. Finally, they developed an approach to derive an empirical Wasserstein radius, which is in spirit similar to our approach in this paper.

1.1 Notation, terminology, and basic techniques

We denote by \mathbb{R}^n the n -dimensional Euclidean space and by \mathbb{R}_+^n the nonnegative orthant in \mathbb{R}^n . For a scalar $p \geq 1$, the p -norm of $z \in \mathbb{R}^n$ is defined $\|z\|_p := (\sum_{i=1}^n |z_i|^p)^{1/p}$, e.g., $\|z\|_1 = \sum_{i=1}^n |z_i|$. We will drop the subscript for the 2-norm, i.e., $\|z\| := \|z\|_2$. For $v, w \in \mathbb{R}^n$, the inner product of v and w is denoted by $v^T w := \sum_{i=1}^n v_i w_i$. For the specific dimensions k and n of the problem in this paper, we denote by e_i the i -th standard basis vector in \mathbb{R}^k , and similarly, denote by f_j the j -th standard basis vector in \mathbb{R}^n . We will also define $g_1 := \begin{pmatrix} e_1 \\ 0 \end{pmatrix} \in \mathbb{R}^{k+n}$. We denote

by δ_ξ the Dirac distribution concentrating unit mass at $\xi \in \mathbb{R}^k$. For any $N \in \mathbb{N}$, we define $[N] := \{1, \dots, N\}$.

Let $\mathbb{R}^{m \times n}$ denote the space of real $m \times n$ matrices, and $A \cdot B := \text{trace}(A^T B)$ denote the trace of the inner product of two matrices $A, B \in \mathbb{R}^{m \times n}$. We denote by \mathbb{S}^n the space of $n \times n$ symmetric matrices, and for $X \in \mathbb{S}^n$, $X \geq 0$ represents that X is positive semidefinite. In addition, we denote by $\text{diag}(X)$ the vector containing the diagonal entries of X , and denote by $\text{Diag}(v)$ the diagonal matrix with vector v along its diagonal. $I \in \mathbb{S}^n$ denotes the identity matrix.

Finally, letting $\mathcal{K} \subseteq \mathbb{R}^n$ be a closed, convex cone, and \mathcal{K}^* be its dual cone, we give a brief introduction to *copositive programming* with respect to the cone \mathcal{K} . The *copositive cone* with respect to \mathcal{K} is defined as

$$\mathcal{COP}(\mathcal{K}) := \{M \in \mathbb{S}^n : x^T M x \geq 0 \forall x \in \mathcal{K}\},$$

and its dual cone, the *completely positive cone* with respect to \mathcal{K} , is given as

$$\mathcal{CP}(\mathcal{K}) := \{X \in \mathbb{S}^n : X = \sum_i x^i (x^i)^T, x^i \in \mathcal{K}\},$$

where the summation over i is finite but its cardinality is unspecified. The term *copositive programming* refers to linear optimization over $\mathcal{COP}(\mathcal{K})$ or, via duality, linear optimization over $\mathcal{CP}(\mathcal{K})$. In fact, these problems are sometimes called *generalized copositive programming* or *set-semidefinite optimization* (Burer and Dong 2012; Eichfelder and Jahn 2008) in contrast with the standard case $\mathcal{K} = \mathbb{R}_+^n$. In this paper, we work with generalized copositive programming, although we use the shorter phrase for convenience.

2 A Wasserstein-based ambiguity set

In this section, we define the Wasserstein metric and discuss a standard method to construct a Wasserstein-based ambiguity set. Using this ambiguity set, we fully specify problem (2).

Denote by $\hat{\Theta}_N := \{\hat{\xi}^1, \dots, \hat{\xi}^N\}$ the set of N independent samples of ξ governed by \mathbb{P} . The uniform empirical distribution based on $\hat{\Theta}_N$ is $\hat{\mathbb{P}}_N := \frac{1}{N} \sum_{i=1}^N \delta_{\hat{\xi}^i}$ where δ_ζ is the Dirac distribution concentrating unit mass at $\zeta \in \mathbb{R}^k$.

Definition 1 (Definition 3 in Hanasusanto and Kuhn (2016)) Let $\mathcal{M}^2(\Xi)$ be the set of all probability distributions \mathbb{Q} that are supported on Ξ and that satisfy $\mathbb{E}_{\mathbb{Q}}[\|\xi - \xi'\|^2] = \int_{\Xi} \|\xi - \xi'\|^2 d\mathbb{Q}(\xi) < \infty$ where $\xi' \in \Xi$ is some reference point, e.g., $\xi' = \hat{\xi}^i$ for some $i \in [N]$.

Definition 2 (Definition 3 in Hanasusanto and Kuhn (2016)) The 2-Wasserstein distance between any $\mathbb{Q}, \mathbb{Q}' \in \mathcal{M}^2(\Xi)$ is

$$W^2(\mathbb{Q}, \mathbb{Q}') := \inf \left\{ \left(\int_{\Xi^2} \|\xi - \xi'\|^2 \Pi(d\xi, d\xi') \right)^{1/2} : \begin{array}{l} \Pi \text{ is a joint distribution of } \xi \text{ and } \xi' \\ \text{with marginals } \mathbb{Q} \text{ and } \mathbb{Q}', \text{ respectively} \end{array} \right\}.$$

Remark 2.1 The Wasserstein distance is essentially the minimum cost of redistributing mass from \mathbb{Q} to \mathbb{Q}' . It is also called the “earth mover’s distance” in the community of computer science; see (Rubner et al. 2000). In fact, the Wasserstein distance between two discrete distributions with a finite number of positive masses corresponds to a transportation planning problem in finite dimensions.

Example 1 illustrates the Wasserstein distance between two discrete distributions.

Example 1 Consider two discrete distributions: $\mathbb{Q} := \sum_{i=1}^M q_i \delta_{\xi_i}$ and $\mathbb{Q}' := \sum_{j=1}^{M'} q'_j \delta_{\xi'_j}$ where $q_i \geq 0$ $i = 1, \dots, M$, $q'_j \geq 0$ $j = 1, \dots, M'$, and $\sum_{i=1}^M q_i = \sum_{j=1}^{M'} q'_j = 1$. Define $c_{ij} = \|\xi_i - \xi'_j\|^2$ $\forall i = 1, \dots, M$ $j = 1, \dots, M'$. Then, the 2-Wasserstein distance between \mathbb{Q} and \mathbb{Q}' equals the square root of the optimal value of the following transportation planning problem:

$$\begin{aligned} & \min_{\pi} \sum_{i=1}^M \sum_{j=1}^{M'} c_{ij} \pi_{ij} \\ & \text{s. t. } \sum_{j=1}^{M'} \pi_{ij} = q_i \quad \forall i = 1, \dots, M \\ & \quad \sum_{i=1}^M \pi_{ij} = q'_j \quad \forall j = 1, \dots, M' \\ & \quad \pi_{ij} \geq 0 \quad \forall i = 1, \dots, M, \quad j = 1, \dots, M', \end{aligned} \quad (3)$$

where π is the joint distribution of ξ and ξ' with marginals of \mathbb{Q} and \mathbb{Q}' and π is the matrix variable in this optimization problem.

With this setting, our ambiguity set contains a family of distributions that are close to $\widehat{\mathbb{P}}_N$ with respect to the Wasserstein metric. In particular, we define our ambiguity set \mathcal{D} as a 2-Wasserstein ball of radius ϵ that is centered at the uniform empirical distribution $\widehat{\mathbb{P}}_N$:

$$\mathcal{D}(\widehat{\mathbb{P}}_N, \epsilon) := \left\{ \mathbb{Q} \in \mathcal{M}^2(\Xi) : W^2(\mathbb{Q}, \widehat{\mathbb{P}}_N) \leq \epsilon \right\}.$$

The reader is referred to (Hanasusanto and Kuhn 2016) for the general case of $\mathcal{M}^r(\Xi)$ and $W^r(\mathbb{Q}, \mathbb{Q}')$ for any $r \geq 1$. We use the 2-Wasserstein distance in this paper for two reasons. First, the Euclidean distance is one of the most popular distances considered in the relevant literature; see (Esfahani and Kuhn 2015; Hanasusanto and Kuhn 2016). Second, we will find that problem (2) with an ambiguity set based on the 2-Wasserstein distance can be reformulated into a copositive program; see Sect. 3.

Then, we replace the generic ambiguity set \mathcal{D} with the Wasserstein ball $\mathcal{D}(\widehat{\mathbb{P}}_N, \epsilon)$ in problem (2) to compute a data-driven upper bound:

$$\begin{aligned} v_{\mathcal{D}(\widehat{\mathbb{P}}_N, \epsilon)}^+ = & \sup_{\Pi, \mathbb{Q} \in \mathcal{M}^2(\Xi)} \int_{\Xi} v(\xi) d\mathbb{Q}(\xi) \\ \text{s. t.} & \int_{\Xi^2} \|\xi - \xi'\|^2 \Pi(d\xi, d\xi') \leq \epsilon^2 \\ & \Pi \text{ is a joint distribution of } \xi \text{ and } \xi' \\ & \text{with marginals } \mathbb{Q} \text{ and } \widehat{\mathbb{P}}_N, \text{ respectively.} \end{aligned} \quad (4)$$

We next close this subsection by making some remarks. First, the Wasserstein ball radius in problem (4) controls the conservatism of the optimal value. A larger radius is more likely to contain the true distribution and thus a more likely valid upper bound on $v_{\mathbb{P}}$, but even if it is valid, it could be a weaker upper bound. Therefore, it is crucial to choose an appropriate radius for the Wasserstein ball. Second, the Kullback–Leibler divergence ball is also considered in recent research; see (Hu and Hong 2013; Jiang and Guan 2016). However, in the case of our discrete empirical distribution, the Kullback–Leibler divergence ball is a singleton containing only the empirical distribution itself, with probability one. Third, the ambiguity sets constructed by goodness-of-fit tests in Bertsimas et al. (2013, 2014) also possess statistical guarantees, however, they often lead to complicated and intractable optimization problems for the case of high-dimensional uncertain parameters.

2.1 An empirical Wasserstein radius

The papers (Eichfelder and Jahn 2008; Esfahani and Kuhn 2015) present a theoretical radius $\epsilon_N(\beta)$ for datasets of size N , which guarantees a desired confidence level $1 - \beta$ for $\mathbb{P} \in \mathcal{D}(\widehat{\mathbb{P}}_N, \epsilon_N(\beta))$ under the following standard assumption on \mathbb{P} :

Assumption 1 (Light-tailed distribution) There exists an exponent $a > 1$ such that

$$\mathbb{E}_{\mathbb{P}}[\exp(\|\xi\|^a)] = \int_{\Xi} \exp(\|\xi\|^a) d\mathbb{P}(\xi) < \infty.$$

Note that $\epsilon_N(\beta)$ depends on N and β . However, $\epsilon_N(\beta)$ is known to be conservative in practice; see (Eichfelder and Jahn 2008) for example. In other words, $\mathcal{D}(\widehat{\mathbb{P}}_N, \epsilon_N(\beta))$ might contain significantly more irrelevant distributions so that the computed $v_{\mathcal{D}(\widehat{\mathbb{P}}_N, \epsilon_N(\beta))}^+$ is significantly larger than $v_{\mathbb{P}}$. So, we propose a procedure to derive an empirical radius that provides a desired confidence level $1 - \beta$ but is much smaller than $\epsilon_N(\beta)$. Our approach is based on the data set $\widehat{\Theta}_N$. In particular, we apply a procedure, similar to cross validation in spirit, that computes an empirical confidence level (between 0 and 1) for a given radius ϵ ; see details in the next paragraphs. Our procedure guarantees that a larger radius leads to a higher confidence level. Therefore, by iteratively testing different ϵ , we can find a radius with a desired confidence level based on the data set $\widehat{\Theta}_N$. Although the derived $\epsilon(\widehat{\Theta}_N, \beta)$ depends on the data set $\widehat{\Theta}_N$, our experimental results in Sect. 4 indicate that it can be used for other datasets of the same sample size. We will show the numerical evidence in Sect. 4. Our approach is

also similar in spirit to the one used in Esfahani and Kuhn (2015); Hanasusanto and Kuhn (2016).

Our procedure requires an oracle to compute (or approximate) $v_{\mathcal{D}(\mathbb{P}_N, \epsilon)}^+$. Later in Sect. 3, we will propose a specific approximation; see (21). Assume also that, in addition to the dataset $\widehat{\Theta}_N$, we predetermine a set \mathcal{E} containing a large, yet finite, number of candidate radii ϵ . We randomly divide $\widehat{\Theta}_N$ into training and validation datasets K times. We enforce the same dataset size denoted by N_T on each of the K training datasets.

Next, for each $\epsilon \in \mathcal{E}$, we derive an empirical probability based on the following procedure: (i) we use each of the K training datasets to approximate $v_{\mathcal{D}(\mathbb{P}_{N_T}, \epsilon)}^+$ with a value called $v_{\text{WB}}(\epsilon)$ by calling the oracle, where \mathbb{P}_{N_T} represents the empirical distribution from the training set; (ii) we then use the corresponding K validation datasets to simulate the expected optimal values denoted by v_{SB} ¹ and (iii) we finally compute the percentage of the K instances where $v_{\text{WB}}(\epsilon) \geq v_{\text{SB}}$. Let us call this empirical probability as the *empirical confidence level*. Thus, the empirical confidence level can roughly approximate the confidence level that the underlying distribution is contained in the Wasserstein-based ambiguity set with the radius ϵ . Note that the percentage computed is non-decreasing in ϵ and equal to 1 for some large ϵ_0 . Therefore, the set containing all the empirical confidence levels is essentially an empirical cumulative distribution. Then, given a desired confidence level, we can choose a corresponding empirical radius $\epsilon \in \mathcal{E}$. The numerical results in Sect. 4 indicate that our choices of ϵ indeed return the desired confidence levels. We specify the above procedure in Algorithm 1.

Algorithm 1 Procedure to compute an empirical confidence level for any $\epsilon \in \mathcal{E}$

Inputs: A dataset $\widehat{\Theta}_N = \{\hat{\xi}^1, \dots, \hat{\xi}^N\}$ and a radius $\epsilon \in \mathcal{E}$

Outputs: The empirical confidence level

for $k = 1, \dots, K$ **do**

 Use the k^{th} training dataset to compute $v_{\text{WB}}^k(\epsilon)$

 Use the k^{th} validation dataset to simulate v_{SB}^k

end for

Calculate the empirical confidence level for ϵ as the percentage of the K instances where $v_{\text{WB}}^k(\epsilon) \geq v_{\text{SB}}^k$

3 Problem reformulation and tractable bound

In this section, we propose a copositive programming reformulation for problem (4) under some mild assumptions. As copositive programs are computationally intractable, we then propose semidefinite-based relaxations for the purposes of computation.

Let us first define the feasible set for $x \in \mathbb{R}^n$ in (1) as follows:

¹ This process is to solve a linear program or integer program corresponding to each sample in the validation dataset and then to take the average of the optimal values.

$$\mathcal{X} := \left\{ x \in \mathbb{R}^n : \begin{array}{l} Ax = b, \ x \geq 0 \\ x_j \in \{0, 1\} \ \forall j \in \mathcal{B} \end{array} \right\}.$$

We now introduce the following standard assumptions:

Assumption 2 The set $\mathcal{X} \subseteq \mathbb{R}^n$ is nonempty and bounded.

Assumption 3 $Ax = b, \ x \geq 0 \implies 0 \leq x_j \leq 1 \ \forall j \in \mathcal{B}$.

Assumption 3 can be easily enforced. For example (see also Burer 2009; Natarajan et al. 2011), if $\mathcal{B} = \emptyset$, then the assumption is redundant; if problem (1) is derived from the network flow problems, for instance the longest path problem on a directed acyclic graph, then Assumption 3 is implied from the network flow constraints; if \mathcal{B} is a nonempty set and the assumption is not implied by the constraints, we can add constraints $x_j + s_j = 1, \ s_j \geq 0 \ \forall j \in \mathcal{B}$.

Assumption 4 The support set $\Xi \subseteq \mathbb{R}^k$ is convex, closed, and computationally tractable.

By tractable, we mean that there exists a polynomial-time algorithm that can determine whether or not a given point lies in Ξ and, if not, can generate a cut (normally a linear inequality) to separate the point; see (Grötschel et al. 1981). For example, Ξ could be represented using a polynomial number of linear, second-order-cone, and semidefinite inequalities.

Assumption 5 Ξ is bounded.

By Assumption 2, we know that $v(\xi)$ is finite and attainable for any $\xi \in \Xi$. Note that under Assumptions 2–5, v_D^+ is finite and attainable and thus we can replace sup with max in (4) under these conditions. Assumption 5 could be merged with Assumption 4, but it is stated separately to highlight its role in proving the exactness of the copositive programming reformulation below.

3.1 A copositive reformulation

We reformulate problem (4) via conic programming duality theory and probability theory. We introduce a useful result from the literature as follows.

Lemma 1 $v_{D(\mathbb{P}_N, \epsilon)}^+$ equals the optimal value of

$$\begin{array}{ll} \sup & \frac{1}{N} \sum_{i=1}^N \int_{\Xi} v(\xi) \, d\mathbb{Q}_i(\xi) \\ \text{s. t.} & \frac{1}{N} \sum_{i=1}^N \int_{\Xi} \|\xi - \hat{\xi}_i\|^2 \, d\mathbb{Q}_i(\xi) \leq \epsilon^2 \\ & \mathbb{Q}_i \in \mathcal{M}^2(\Xi) \ \forall i \in [N], \end{array} \quad (5)$$

where \mathbb{Q}_i represents the distribution of ξ conditional on $\xi' = \hat{\xi}^i$ for all $i \in [N]$.

Proof As \mathbb{Q}_i represents the distribution of ξ conditional on $\xi' = \hat{\xi}^i$, the joint probability Π in problem (4) can be decomposed as $\Pi = \frac{1}{N} \sum_{i \in [N]} \mathbb{Q}_i$ by the law of total probability. Thus, the optimal value of (5) coincides with $v_{\mathcal{D}(\widehat{\mathbb{P}}_N, \epsilon)}^+$, which completes the proof. \square

We next provide a copositive programming reformulation for problem (5). As the first step, we use a standard duality argument to write the dual of (5) (see also Esfahani and Kuhn 2015):

$$v_{\mathcal{D}(\widehat{\mathbb{P}}_N, \epsilon)}^+ = \sup_{\mathbb{Q}_i \in \mathcal{M}^2(\Xi)} \inf_{\lambda \geq 0} \frac{1}{N} \sum_{i=1}^N \int_{\Xi} v(\xi) d\mathbb{Q}_i(\xi) + \lambda \left(\epsilon^2 - \frac{1}{N} \sum_{i=1}^N \int_{\Xi} \|\xi - \hat{\xi}_i\|^2 d\mathbb{Q}_i(\xi) \right) \quad (6)$$

$$\leq \inf_{\lambda \geq 0} \sup_{\mathbb{Q}_i \in \mathcal{M}^2(\Xi)} \lambda \epsilon^2 + \frac{1}{N} \sum_{i=1}^N \int_{\Xi} (v(\xi) - \lambda \|\xi - \hat{\xi}_i\|^2) d\mathbb{Q}_i(\xi) \quad (7)$$

$$= \inf_{\lambda \geq 0} \lambda \epsilon^2 + \frac{1}{N} \sum_{i=1}^N \sup_{\xi \in \Xi} (v(\xi) - \lambda \|\xi - \hat{\xi}_i\|^2), \quad (8)$$

where (7) follows from the max–min inequality, while Eq. (8) follows from the fact that $\mathcal{M}^2(\Xi)$ contains all the Dirac distributions supported on Ξ .

By Assumption 2, $v(\xi)$ is finite for all $\xi \in \Xi$. Then, the inequality in (7) becomes an equality for any $\epsilon > 0$ due to a straightforward generalization of a strong duality result for moment problems in Proposition 3.4 in Shapiro (2001); see also Theorem 1 in Hanasusanto and Kuhn (2016) and Lemma 7 in Hanasusanto et al. (2017). By introducing auxiliary variables s_i , the minimization problem in (8) is equivalent to

$$v_{\mathcal{D}(\widehat{\mathbb{P}}_N, \epsilon)}^+ = \inf_{\lambda, s_i} \lambda \epsilon^2 + \frac{1}{N} \sum_{i=1}^N s_i \quad \text{s. t. } \sup_{\xi \in \Xi} (v(\xi) - \lambda \|\xi - \hat{\xi}_i\|^2) \leq s_i \quad \forall i \in [N] \quad (9) \\ \lambda \geq 0.$$

For each $i \in [N]$, consider the following maximization problem corresponding to the left-hand side of the constraints in (9):

$$h^i(\lambda) := \sup (F\xi)^T x - \lambda \left(\xi^T \xi - 2\hat{\xi}_i^T \xi + \|\hat{\xi}_i\|^2 \right) \quad \text{s. t. } Ax = b, \quad x \geq 0 \quad (10) \\ x_j \in \{0, 1\} \quad \forall j \in \mathcal{B} \\ e_1^T \xi = 1, \quad \xi \in \hat{\Xi},$$

which is a mixed 0-1 bilinear program. Under Assumption 3, it holds also that the optimal value of (10) equals the optimal value of an associated copositive program (Burer 2009, 2012), which we now describe.

Define

$$z := \begin{pmatrix} \xi \\ x \end{pmatrix} \in \mathbb{R}^{k+n}, \quad E := (-be_1^T \ A) \in \mathbb{R}^{m \times (k+n)}, \quad (11)$$

$$H^i(\lambda) := \begin{pmatrix} -\lambda \left(I - \hat{\xi}_i e_1^T - e_1 \hat{\xi}_i^T + \|\hat{\xi}_i\|^2 e_1 e_1^T \right) \frac{1}{2} F^T \\ \frac{1}{2} F \\ 0 \end{pmatrix} \in \mathbb{S}^{k+n}, \quad (12)$$

and for any $j \in \mathcal{B}$, define

$$Q_j := \begin{pmatrix} 0 \\ f_j \end{pmatrix} \begin{pmatrix} 0 \\ f_j \end{pmatrix}^T - \frac{1}{2} \begin{pmatrix} 0 \\ f_j \end{pmatrix} \begin{pmatrix} e_1 \\ 0 \end{pmatrix}^T - \frac{1}{2} \begin{pmatrix} e_1 \\ 0 \end{pmatrix} \begin{pmatrix} 0 \\ f_j \end{pmatrix}^T \in \mathbb{S}^{k+n}. \quad (13)$$

where f_j denotes the j -th standard basis vector in \mathbb{R}^n .

Because both \mathcal{X} and Ξ are bounded by Assumptions 2 and 5, there exists a scalar $r > 0$ such that the constraint $z^T z = \xi^T \xi + x^T x \leq r$ is redundant for (10). Furthermore, it is well-known that we can use the following quadratic constraints to represent the binary variables in the description of \mathcal{X} :

$$x_j^2 - x_j = 0 \quad \Leftrightarrow \quad Q_j \cdot z z^T = 0.$$

After adding the redundant constraint and representing the binary variables, we homogenize problem (10) as follows:

$$\begin{aligned} & \max \quad H^i(\lambda) \bullet z z^T \\ & \text{s. t.} \quad E z = 0, \quad g_1^T z = 1 \\ & \quad \quad I \cdot z z^T \leq r \\ & \quad \quad Q_j \cdot z z^T = 0 \quad \forall j \in \mathcal{B} \\ & \quad \quad z \in \widehat{\Xi} \times \mathbb{R}_+^n, \end{aligned} \quad (14)$$

where $g_1 = \begin{pmatrix} e_1 \\ 0 \end{pmatrix} \in \mathbb{R}^{k+n}$ and e_1 denotes the standard basis vector in \mathbb{R}^k . The copositive representation is thus

$$\begin{aligned} & \max \quad H^i(\lambda) \cdot Z \\ & \text{s. t.} \quad \text{diag}(E Z E^T) = 0 \\ & \quad \quad g_1 g_1^T \cdot Z = 1 \\ & \quad \quad I \cdot Z \leq r \\ & \quad \quad Q_j \cdot Z = 0 \quad \forall j \in \mathcal{B} \\ & \quad \quad Z \in \mathcal{CP}(\widehat{\Xi} \times \mathbb{R}_+^n). \end{aligned} \quad (15)$$

Letting $u^i \in \mathbb{R}^m$, $\rho^i \in \mathbb{R}_+$, $\alpha^i \in \mathbb{R}$, and $v^i \in \mathbb{R}^{|\mathcal{B}|}$ be the respective dual multipliers of $\text{diag}(E Z E^T) = 0$, $I \cdot Z \leq r$, $g_1 g_1^T \cdot Z = 1$, and $Q_j \cdot Z = 0$, standard conic duality theory implies the dual of (15) is

$$\begin{aligned}
& \min_{\alpha^i, \rho^i, u^i, v^i} \alpha^i + r\rho^i \\
& \text{s. t.} \quad \alpha^i g_1 g_1^T - H^i(\lambda) + E^T \text{Diag}(u^i)E + \sum_{j \in \mathcal{B}} v_j^i Q_j + \rho^i I \in \mathcal{COP}(\widehat{\mathcal{E}} \times \mathbb{R}_+^n) \\
& \quad \rho^i \geq 0.
\end{aligned} \tag{16}$$

Holding all other dual variables fixed, for $\rho^i > 0$ large, the matrix variable in (16) is strictly copositive—in fact, positive definite—which establishes that Slater's condition is satisfied, thus ensuring strong duality: the optimal value of (15) equals the optimal value of (16). Therefore, we can reformulate problem (9) as follows:

$$\begin{aligned}
v_{\mathcal{D}(\widehat{\mathbb{P}}_N, \epsilon)}^+ &= \min \lambda \epsilon^2 + \frac{1}{N} \sum_{i=1}^N (\alpha^i + r\rho^i) \\
& \text{s. t.} \quad \alpha^i g_1 g_1^T - H^i(\lambda) + E^T \text{Diag}(u^i)E + \sum_{j \in \mathcal{B}} v_j^i Q_j + \rho^i I \in \mathcal{COP}(\widehat{\mathcal{E}} \times \mathbb{R}_+^n) \quad \forall i \in [N] \\
& \quad \rho^i \geq 0 \quad \forall i \in [N] \\
& \quad \lambda \geq 0.
\end{aligned} \tag{17}$$

Note that if Assumption 5 fails, the constraint $I \cdot Z \leq r$ should be excluded from (15) and thus the terms $r\rho^i$ and $\rho^i I$ in the objective function and the constraint, respectively, should be excluded in (16) as well. As such, strong duality between (15) and (16) cannot be established in this case. However, the modified (16) still provides an upper bound on $h^i(\lambda)$. Accordingly, the modified problem (17) still provides an upper bound on $v_{\mathcal{D}(\widehat{\mathbb{P}}_N, \epsilon)}^+$.

3.2 A semidefinite-based relaxation

As problem (17) is difficult to solve in general, we propose a tractable approximation based on semidefinite programming techniques. In particular, we propose an inner approximation of $\mathcal{COP}(\widehat{\mathcal{E}} \times \mathbb{R}_+^n)$ in (17) so that the resulting problem has an optimal value that is an upper bound on $v_{\mathcal{D}}^+$. Now, define

$$\text{IA}(\widehat{\mathcal{E}} \times \mathbb{R}_+^n) := \left\{ S + M : \begin{array}{l} S_{11} \in \text{IA}(\widehat{\mathcal{E}}), \text{Rows}(S_{21}) \in \widehat{\mathcal{E}}^* \\ S_{22} \geq 0, M \succeq 0 \end{array} \right\},$$

where $\text{IA}(\widehat{\mathcal{E}})$ is an inner approximation of $\mathcal{COP}(\widehat{\mathcal{E}})$, i.e., $\text{IA}(\widehat{\mathcal{E}}) \subseteq \mathcal{COP}(\widehat{\mathcal{E}})$. Immediately, we have a relationship between $\text{IA}(\widehat{\mathcal{E}} \times \mathbb{R}_+^n)$ and $\mathcal{COP}(\widehat{\mathcal{E}} \times \mathbb{R}_+^n)$:

Lemma 2 $\text{IA}(\widehat{\mathcal{E}} \times \mathbb{R}_+^n) \subseteq \mathcal{COP}(\widehat{\mathcal{E}} \times \mathbb{R}_+^n)$.

Proof Let arbitrary $\begin{pmatrix} p \\ q \end{pmatrix} \in \widehat{\mathcal{E}} \times \mathbb{R}_+^n$ be given. We need to show

$$\begin{aligned}
\begin{pmatrix} p \\ q \end{pmatrix}^T (S + M) \begin{pmatrix} p \\ q \end{pmatrix} &= \begin{pmatrix} p \\ q \end{pmatrix}^T S \begin{pmatrix} p \\ q \end{pmatrix} + \begin{pmatrix} p \\ q \end{pmatrix}^T M \begin{pmatrix} p \\ q \end{pmatrix} \geq 0. \\
\begin{pmatrix} p \\ q \end{pmatrix}^T (S + M) \begin{pmatrix} p \\ q \end{pmatrix} &= \begin{pmatrix} p \\ q \end{pmatrix}^T S \begin{pmatrix} p \\ q \end{pmatrix} + \begin{pmatrix} p \\ q \end{pmatrix}^T M \begin{pmatrix} p \\ q \end{pmatrix}
\end{aligned} \tag{18}$$

$$= p^T S_{11} p + 2q^T S_{21} p + q^T S_{22} q + \binom{p}{q}^T M \binom{p}{q} \quad (19)$$

$$\geq 0 \quad (20)$$

The first term is nonnegative because $p \in \widehat{\Xi}$ and $S_{11} \in \text{IA}(\widehat{\Xi}) \subseteq \mathcal{COP}(\widehat{\Xi})$; the second term is nonnegative because $p \in \widehat{\Xi}$, $q \geq 0$, and $\text{Rows}(S_{21}) \in \widehat{\Xi}^*$; the third term is nonnegative because $q \geq 0$ and $S_{22} \geq 0$; the last term is nonnegative because $M \succeq 0$.

When $\widehat{\Xi} = \{\xi \in \mathbb{R}^k : P\xi \geq 0\}$ is a polyhedral cone based on some matrix $P \in \mathbb{R}^{p \times k}$, a typical inner approximation $\text{IA}(\widehat{\Xi})$ of $\mathcal{COP}(\widehat{\Xi})$ is given by

$$\text{IA}(\widehat{\Xi}) := \{S_{11} = P^T Y P : Y \geq 0\},$$

where $Y \in \mathbb{S}^p$ is a symmetric matrix variable. This corresponds to the RLT approach of (Anstreicher 2009; Burer 2015; Serali and Adams 2013). When $\widehat{\Xi} = \{\xi \in \mathbb{R}^k : \|(\xi_2, \dots, \xi_k)^T\| \leq \xi_1\}$ is the second-order cone, it is known (Sturm and Zhang 2003) that

$$\mathcal{COP}(\widehat{\Xi}) = \{S_{11} = \tau J + M_{11} : \tau \geq 0, M_{11} \geq 0\},$$

where $J = \text{Diag}(1, -1, \dots, -1)$. Because of this simple structure, it often makes sense to take $\text{IA}(\widehat{\Xi}) = \mathcal{COP}(\widehat{\Xi})$ in practice.

Now consider the following problem by replacing $\mathcal{COP}(\widehat{\Xi} \times \mathbb{R}_+^n)$ with $\text{IA}(\widehat{\Xi} \times \mathbb{R}_+^n)$ in (17).

$$\begin{aligned} \bar{v}_{\mathcal{D}(\widehat{\mathbb{P}}, \epsilon)}^+ &= \min \lambda \epsilon^2 + \frac{1}{N} \sum_{i=1}^N (\alpha^i + r \rho^i) \\ \text{s.t. } &\alpha^i g_{181}^T - H^i(\lambda) + E^T \text{Diag}(u^i) E + \sum_{j \in \mathcal{B}} v_j^i Q_j + \rho^i I \in \text{IA}(\widehat{\Xi} \times \mathbb{R}_+^n) \quad \forall i \in [N] \\ &\rho^i \geq 0 \quad \forall i \in [N] \\ &\lambda \geq 0. \end{aligned} \quad (21)$$

Obviously, we have the following result:

Theorem 1 $v_{\mathcal{D}(\widehat{\mathbb{P}}_N, \epsilon)}^+ \leq \bar{v}_{\mathcal{D}(\widehat{\mathbb{P}}, \epsilon)}^+.$

4 Numerical experiments

In this section, we validate our proposed Wasserstein-ball approach (WB) on three applications. We will compare WB with the moment-based approach (MB) proposed in Natarajan et al. (2011) where the exact values of the first two moments of the distributions are known. In practice, the moments of the distribution are often not known exactly. To this end, Delage and Ye (2010) proposed a data-driven approach to handle this case. However, in this paper, we assume that the moments are known exactly for MB. Actually, this choice favors MB, but the goal of our experiments is to demonstrate that our approach provides a valid upper bound that gets closer to $v_{\mathbb{P}}$ as the size of the data set increases, while the MB provides an upper bound, which does not improve with the size of the data set.

All computations are conducted with Mosek version 8.0.0.28 beta (MOSEK ApS 2016) on an Intel Core i3 2.93 GHz Windows computer with 4 GB of RAM and are implemented using the modeling language YALMIP (Lofberg 2004) in MATLAB (R2014a) version 8.3.0.532. In order to demonstrate the effectiveness of WB, we also implement a Monte Carlo simulation-based approach (SB) which requires a sufficiently large number of randomly generated samples. For the project management problem in Sect. 4.2, a linear program is solved for each sample of the Monte Carlo simulation, while for the knapsack problem in Sect. 4.3, an integer program is solved for each sample. We employ CPLEX 12.4 to solve these linear programs and integer programs.

4.1 Statistical sensitivity analysis of highest-order statistic

The problem of finding the maximum value from a set $\zeta = (\zeta_1, \dots, \zeta_n)$ of n numbers can be formulated as the optimization problem:

$$\max \left\{ \zeta^T x : e^T x = 1, x \geq 0 \right\}. \quad (22)$$

For example, suppose $\zeta_1 = \max\{\zeta_1, \dots, \zeta_n\}$, then the optimal solution to (22) is $x_1^* = 1, x_2^* = \dots = x_n^* = 0$. For the statistical sensitivity analysis problem, we consider a random vector ζ following a joint distribution \mathbb{P} . In the situation where the true distribution is not known exactly, our focus is to investigate the upper bound on the expected maximum value over an ambiguity set containing distributions that possess partial shared information.


We consider an instance with $n = 3$ and the true distribution \mathbb{P} of ζ is assumed to be jointly lognormal with first and second moments given by $\mu_{\log} \in \mathbb{R}^3$ and $\Sigma_{\log} \in \mathbb{S}^3$, respectively.

In our experiments, we use the following procedure to randomly generate μ_{\log} and Σ_{\log} . We first sample $\mu \in \mathbb{R}^3$ from a uniform distribution $[0, 2]^3$. Then, we randomly generate a matrix $\Sigma \in \mathbb{S}^3$ as follows: we set the vector of standard deviations to $\sigma = \frac{1}{4}e \in \mathbb{R}^3$, sample a random correlation matrix $C \in \mathbb{S}^3$ using the MATLAB command ‘gallery(‘randcorr’,3)’, and set $\Sigma = \text{diag}(\sigma)C \text{diag}(\sigma) + \mu\mu^T$. We set μ and Σ as the first and second moments respectively of the corresponding normal distribution of \mathbb{P} . Then μ_{\log} and Σ_{\log} can be computed based on the following formulae (Halliwell 2015):

$$\begin{aligned} (\mu_{\log})_i &= e^{\mu_i + 0.5\Sigma_{ii}}, \\ (\Sigma_{\log})_{ij} &= e^{\mu_i + \mu_j + 0.5(\Sigma_{ii} + \Sigma_{jj})} (e^{\Sigma_{ij}} - 1). \end{aligned} \quad (23)$$

We can cast this problem into our framework by setting $m = 1, k = n + 1$, $\xi = (1, \zeta_1, \dots, \zeta_n)$, $F = (0, I)$, and $\mathcal{B} = \emptyset$. Obviously, Assumptions 2 and 4 are satisfied. Assumption 3 is vacuous. Although Assumption 5 does not hold, problem (21) can still provide a valid upper bound on the expected optimal value as discussed in Sect. 3.1.

Fig. 1 Empirical confidence levels of different Wasserstein radii for $N \in \{20, 80, 320, 1280\}$ respectively



4.1.1 The deviation of empirical Wasserstein radii


In this experiment, we consider a particular underlying distribution \mathbb{P} that is generated by the procedure mentioned above. Also, we consider eight cases for the size of the dataset: $N \in \{10, 20, 40, 80, 160, 320, 640, 1280\}$. For each case, we randomly generate a dataset $\hat{\Theta}_N$ containing N independent samples from \mathbb{P} and use the procedure in Sect. 2 to determine a desired radius from a pre-specified set $\mathcal{E} = \{0.001, 0.005, 0.01, 0.05, 0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1.0, 2.0\}$.² In particular, we set $K = 100$ in Algorithm 1. Figure 1 shows the trend of the reliabilities over different Wasserstein radii for $N \in \{20, 80, 320, 1280\}$. Clearly, smaller Wasserstein radii tend to have lower empirical confidence levels. Furthermore, as the sample size increases, the empirical confidence level increases as well for the same Wasserstein radius. The result of this experiment indicates that we can practically choose a Wasserstein radius with a desired statistical guarantee for each case of N . We remark that the derived radii can be used for datasets of the same sizes generated from different distributions of the same family (lognormal distributions in this application).

4.1.2 Instances with the same underlying distribution

Our next experiment is to focus on a particular joint lognormal distribution \mathbb{P} . We consider eight cases: $N \in \{10, 20, 40, 80, 160, 320, 640, 1280\}$. For each case, we test 100 trials and in each trial we randomly generate N independent samples from \mathbb{P} and choose the Wasserstein radius with an empirical confidence level of 0.90. We compare our approach with MB where the first two moments are directly given as μ_{\log} and Σ_{\log} . We also randomly generate 100,000 independent samples from \mathbb{P} to simulate the true expected optimal value.

² From preliminary experiments, the largest element 2.0 in set \mathcal{E} returned 1 as the empirical confidence level for all the experiments we conducted. Thus, we believe it is sufficient to have 2.0 as the largest element here.

Fig. 2 The comparison of WB and MB for the stochastic sensitivity analysis problem over different sample sizes for a particular randomly generated underlying distribution



We demonstrate experimental results in Fig. 2. Note that the solid black line represents the simulated value of the true expected optimal value, while the dashed black line represents the upper bound calculated by the moment-based approach. Furthermore, we solve an instance of (21) for each of the 100 trials in each case of N . We use the blue, red, and green lines to respectively represent the 80th quantile, the median, and the 20th quantile of the values from the 100 trials in each case. Figure 2 shows that our approach provides weaker bounds on the expected optimal value for smaller sample sizes. However, as the size of samples increases, our approach provides stronger bounds and the bounds get relatively close to the simulated value. In contrast, the value from MB remains the same regardless of sample size.

4.1.3 Instances with different underlying distributions

In this experiment, we consider eight cases $N \in \{10, 20, 40, 80, 160, 320, 640, 1280\}$. In each case, we randomly generate 100 trials. For each trial in each case we generate N samples from a random lognormal distribution whose first and second moments are generated by using the procedure at the beginning of this section. For each trial in each case, we solve an instance of (21) with a Wasserstein radius corresponding to an empirical confidence level of 0.90. We also simulate the true expected optimal values by randomly generating 100,000 samples from the true distributions.

For each trial in each case, we denote the optimal value from (21) by \bar{v}_{WB}^+ and the simulated value by v_{SB} . Then, we calculate the relative gap between WB and SB as

$$\text{gap}(\text{WB}) := \frac{\bar{v}_{WB}^+ - v_{SB}}{v_{SB}}.$$

We take the average of the relative gaps over the 100 trials for each case. Then, for each trial in each case, we solve MB with the first two moments computed by (23). Denote the optimal value from MB by \bar{v}_{MB}^+ . Similarly, we calculate the relative gap between MB and SB as

Fig. 3 The average gaps from MB and WB for the stochastic sensitivity analysis problem over the eight cases: $N \in \{10, 20, 40, 80, 160, 320, 640, 1280\}$. The blue line represents the average relative gap between the optimal value from WB and the simulated value; the red line represents the average relative gap between the optimal value from MB and the simulated value




Table 1 The percentage of the 100 trials where the optimal values from WB are greater or equal to the simulated values over the 8 cases for the stochastic sensitivity analysis problem

Case number	1	2	3	4	5	6	7	8
Empirical confidence level	1.00	1.00	1.00	1.00	0.97	0.98	0.97	0.98

$$\text{gap}(\text{MB}) := \frac{\bar{v}_{\text{MB}}^+ - v_{\text{SB}}}{v_{\text{SB}}}.$$


We then take the average of the relative gaps over the 100 trials in each case. Figure 3 illustrates the average relative gaps from both WB and MB over the eight cases. Clearly, the upper bound from WB approaches the simulated value along with the increase of the size of samples, while the average relative gap between the bound from MB and the simulated value does not.

Table 1 shows the percentage of the 100 trials where the optimal values from WB are greater than or equal to the corresponding simulated optimal values in the eight cases. The result demonstrates that the derived empirical Wasserstein radii indeed provide desired statistical guarantees in practice.

4.2 Project management problem

In this application, we consider a project management problem, which can be formulated as a longest-path problem on a directed acyclic graph. The arcs denote activities and the nodes denote the completions of a set of activities. Arc lengths denote the time to complete the activities. Thus, the longest path from the starting node s to the ending node t gives the time needed to complete the whole project. Let ζ_{ij} be the length (time) of arc (activity) from node i to node j . The problem can be solved as a linear program due to the network flow structure as follows:

Fig. 4 The structure of the project network where 1 and 6 are the starting and ending nodes respectively



$$\begin{aligned}
 & \max \sum_{(i,j) \in \mathcal{A}} \zeta_{ij} x_{ij} \\
 & \text{s. t.} \quad \sum_{i:(i,j) \in \mathcal{A}} x_{ij} - \sum_{j:(i,j) \in \mathcal{A}} x_{ji} = \begin{cases} 1, & \text{if } i = s \\ 0, & \text{if } i \in \mathcal{N}, \text{ and } i \neq s, t \\ -1, & \text{if } i = t \end{cases} \quad (24) \\
 & \quad x_{ij} \geq 0, \forall (i, j) \in \mathcal{A},
 \end{aligned}$$

where \mathcal{A} denotes the set containing all the arcs, \mathcal{N} denotes the set containing all nodes on the network, and x_{ij} denotes the number of units of flow sent from node i to node j through arc $(i, j) \in \mathcal{A}$. For the stochastic project management problem, the activity times are random. In such cases, due to the resource allocation and management constraints, the project manager would like to quantify the worst-case expected completion time of the project, which is corresponding to the worst-case longest path of the network.

We consider an instance with a network structure shown in Fig. 4. This network consists of 7 arcs and 6 nodes. There are 3 paths from the starting node to the ending node on the network. In the experiments of this example, we consider truncated joint normal distributions. We use the following procedure to generate a truncated joint normal distribution \mathbb{P} : denoting $|\mathcal{A}|$ by the cardinality of set \mathcal{A} , we generate $\zeta \geq 0$ from a jointly normal distribution with first and second moments given by $\mu \in \mathbb{R}^{|\mathcal{A}|}$ and $\Sigma \in \mathbb{S}^{|\mathcal{A}|}$, respectively. Specifically, we sample μ from a uniform distribution $[0, 5]^{|\mathcal{A}|}$ while the matrix Σ is generated randomly using the following procedure: we set the vector of standard deviations to $\sigma = e$, sample a random correlation matrix $C \in \mathbb{S}^{|\mathcal{A}|}$ using the MATLAB command ‘gallery(‘randcorr’, $|\mathcal{A}|$)’, and set $\Sigma = \text{diag}(\sigma)C \text{diag}(\sigma) + \mu\mu^T$. Skipping the details, we can cast the network flow problem into our framework. It is straightforward to check that Assumptions 2, 4, and 5 are satisfied and Assumption 3 is vacuous.

4.2.1 Instances with the same underlying distribution

The first experiment of this example focuses on a particular underlying distribution \mathbb{P} . We consider seven cases: $N \in \{10, 20, 40, 80, 160, 320, 640\}$. For each case, we run 100 trials and in each trial we randomly generate a dataset $\hat{\Theta}_N$ containing N independent samples from \mathbb{P} . We use the procedure in Sect. 2 to compute an empirical confidence level set for each case. Then, we use computed empirical confidence level sets to derive empirical Wasserstein radii for the following computations. For each trial

Fig. 5 The comparison of WB and MB for the project management problem over different sample sizes for a particular randomly generated underlying distribution





Fig. 6 The average gaps from both MB and WB for the project management problem over the seven cases: $N \in \{10, 20, 40, 80, 160, 320, 640\}$



in each case, we solve an instance of (21) with a Wasserstein radius corresponding to an empirical confidence level of 0.90. We compare WB with MB where the first two moments are approximated by using the sample mean and variance from 100,000 samples. The computed moments are close to their theoretical counterparts as the sample size is considerably large. We also simulate the expected optimal value over the 100,000 samples. Figure 5 shows that WB provides weaker bounds on the expected optimal value for smaller sample sizes. However, as the size of samples increases, WB provides stronger bounds and the bounds get relatively close to the simulated value. In contrast, the bounds from MB remains the same regardless of the change of sample sizes.

4.2.2 Instances with different underlying distributions

In this experiment, we consider seven cases: $N \in \{10, 20, 40, 80, 160, 320, 640\}$. For each case, we randomly generate 100 trials in which N independent samples are drawn

Table 2 The percentage of the 100 trials where the optimal values from WB are greater than or equal to the corresponding simulated optimal values over the 7 cases for the project management problem

Case number	1	2	3	4	5	6	7
Empirical confidence level	1.00	1.00	1.00	1.00	0.97	0.98	0.97

from a randomly generated truncated joint normal distribution. Then, for each trial in each case, we solve an instance of (21) with a Wasserstein radius corresponding to a 0.90 empirical confidence level. We solve MB where the first two moments are approximated by computing the sample mean and variance of 100,000 samples. We also simulate the expected optimal value over the 100,000 samples for each trial in each case. We compute the relative gap between the WB and SB as well as the relative gap between MB and SB. Then, for each case, we take the average of the relative gaps from both WB and MB over the 100 trials. Figure 6 illustrates the average relative gaps over the seven cases. Clearly, the upper bound from WB approaches to the simulated value along with the increase in the size of samples, while the gap between the bound from MB and the simulated value remains relatively the same as the sample size increases. Table 2 shows the percentage of the 100 trials where the optimal values from WB are greater than or equal to the corresponding simulated optimal values over the seven cases.

4.3 Knapsack problem


A standard knapsack problem is defined as follows: given a set of items and each with a weight and a value, the problem is to determine the number of items to include in a knapsack such that the total weight is less than or equal to a given capacity limit and the total value is maximized; see the detail in Martello and Toth (1990). Let w_i and ζ_i be the weight and value of item i ($i = 1, \dots, n$) respectively. Let W be the maximum weight capacity of the knapsack. Then, the knapsack problem can be formulated as an integer program:

$$v(\zeta) := \max \left\{ \sum_{i=1}^n \zeta_i x_i : \sum_{i=1}^n w_i x_i \leq W, x_i \in \{0, 1\} \right\},$$

where x_i represents the number of item i to include in the knapsack. Assume that the values of the items are random and follow an unknown joint distribution. Assume also that we can collect a dataset containing N samples with each corresponding to an observation of the n item values. In such cases, we would like to compute a data-driven distributionally robust upper bound on the expected maximum value of the knapsack. We can approximate the upper bound by solving problem (21).

We consider an instance with $n = 4$, $w = (5, 4, 6, 3)^T$, and $W = 10$. The true distribution \mathbb{P} of ζ is assumed to be jointly lognormal with first and second moments given by $\mu_{\log} \in \mathbb{R}^4$ and $\Sigma_{\log} \in \mathbb{S}^4$, respectively. Similar to the procedure described in Sect. 4.1, we sample $\mu \in \mathbb{R}^4$ from a uniform distribution $[0, 2]^4$. Then, we randomly

Fig. 7 The comparison of WB and MB for the knapsack problem over different sample sizes for a particular randomly generated underlying distribution



generate a matrix $\Sigma \in \mathbb{S}^4$ as follows: we set the vector of standard deviations to $\sigma = \frac{1}{4}e \in \mathbb{R}^4$, sample a random correlation matrix $C \in \mathbb{S}^4$ using the MATLAB command ‘gallery(‘randcorr’,4)’, and set $\Sigma = \text{diag}(\sigma)C \text{diag}(\sigma) + \mu\mu^T$. Then μ_{\log} and Σ_{\log} can be computed based on (23). We can easily cast this problem into our framework. For simplicity, we skip the details. It is also straightforward to check that the conditions in Assumptions 2, 3, and 4 are satisfied. Although Assumption 5 is not satisfied, we still can solve (21) to obtain a valid upper bound on the expected optimal value of the knapsack problem.

4.3.1 Instances with the same underlying distribution

In the first experiment, we focus on a particular underlying distribution \mathbb{P} and consider eight cases: $N \in \{10, 20, 40, 80, 160, 320, 640, 1280\}$. For each case, we run 100 trials and in each trial we randomly generate a dataset $\hat{\Theta}_N$ containing N independent samples from \mathbb{P} . Similarly, we derive empirical Wasserstein radii for each case. Then, for each trial in each case, we solve an instance of (21) with an empirical Wasserstein radius corresponding to an empirical confidence level of 0.90. We compare our approach with MB where the first two moments are computed by (23). We simulate the expected optimal value over 100,000 samples. Note that we solve an integer program for each sample in the simulation. Figure 7 shows that WB provides weaker bounds on the expected optimal value for smaller sample sizes. However, as the size of samples increases, WB provides stronger bounds and the bounds get relatively close to the simulated value. In contrast, the bounds from MB remains the same regardless of the change of sample sizes. We remark that the upper bound computed by WB may not converge to the true expected optimal value as the sample size increases to infinity; see the trend shown in Fig. 7. This is due to the fact that problem (21) is a relaxation of problem (17) and the fact the relaxation is not tight in this example.

Fig. 8 Illustration of the average gaps from both MB and WB in the case of $N \in \{10, 20, 40, 80, 160, 320, 640\}$. The blue line represents the average gap between the optimal values from WB and the simulated values; the red line represents the average gap between the optimal values from MB and the simulated values




Table 3 The percentage of the 100 trials where the optimal values from WB are greater than or equal to the corresponding simulated optimal values over the 8 cases for the knapsack problem

Case number	1	2	3	4	5	6	7	8
Empirical confidence level	1.00	1.00	1.00	1.00	1.00	1.00	1.00	1.00

4.3.2 Instances with different underlying distributions

This experiment considers eight cases: $N \in \{10, 20, 40, 80, 160, 320, 640, 1280\}$. For each case, we randomly generate 100 trials with each trial is drawn from a randomly generated joint lognormal distribution. Then, for each trial in each case, we solve an instance of (21) with a Wasserstein radius corresponding to a 0.90 empirical confidence level. We simulate the expected optimal value over 100,000 samples for each trial in each case. Next, we compute the relative gap between the values of WB and SB as well as the relative gap between the values of MB and SB. Then, for each case, we take the average of the relative gaps from both WB and MB over the 100 trials. Figure 8 illustrates the average relative gaps over the seven cases. Clearly, the WB gap becomes narrower as the sample size increases, while the MB gap remains relatively the same.

Table 3 shows the percentage of the 100 trials where the optimal values from WB are greater than or equal to the corresponding simulated optimal values over the eight cases.

5 Concluding remarks

In this paper, we have studied the expected optimal value of a mixed 0-1 programming problem with uncertain objective coefficients following a joint distribution whose information is not known exactly but a set of independent samples can be collected. Using the samples, we have constructed a Wasserstein-based ambiguity set that con-

tains the true distribution with a desired confidence level. We proposed an approach to compute the upper bound on the expected optimal value. Then under mild assumption, the problem was reformulated to a copositive program, which leads to a semidefinite-based relaxation. We have validated the effectiveness of our approach over three applications.

Acknowledgements The authors sincerely thank Kurt Anstreicher, Qihang Lin, Luis Zuluaga, and Tianbao Yang for many useful suggestions, which help us improve the results of the paper. The authors are sincerely grateful to anonymous referees for their comments that have improved the paper.

References

- Aldous DJ (2001) The $\zeta(2)$ limit in the random assignment problem. *Random Struct Algorithms* 18(4):381–418
- Anstreicher KM (2009) Semidefinite programming versus the reformulation-linearization technique for nonconvex quadratically constrained quadratic programming. *J Glob Optim* 43(2–3):471–484
- Ben-Tal A, Den Hertog D, De Waegenaere A, Melenberg B, Rennen G (2013) Robust solutions of optimization problems affected by uncertain probabilities. *Manag Sci* 59(2):341–357
- Bertsimas D, Popescu I (2005) Optimal inequalities in probability theory: a convex optimization approach. *SIAM J Optim* 15(3):780–804
- Bertsimas D, Natarajan K, Teo C-P (2004) Probabilistic combinatorial optimization: moments, semidefinite programming, and asymptotic bounds. *SIAM J Optim* 15(1):185–209
- Bertsimas D, Natarajan K, Teo C-P (2006) Persistence in discrete optimization under data uncertainty. *Math Program* 108(2):251–274
- Bertsimas D, Gupta V, Kallus N (2013) Data-driven robust optimization. arXiv preprint [arXiv:1401.0212](https://arxiv.org/abs/1401.0212)
- Bertsimas D, Gupta V, Kallus N (2014) Robust SAA. arXiv preprint [arXiv:1408.4445](https://arxiv.org/abs/1408.4445)
- Bowman RA (1995) Efficient estimation of arc criticalities in stochastic activity networks. *Manag Sci* 41(1):58–67
- Burer S (2009) On the copositive representation of binary and continuous nonconvex quadratic programs. *Math Program* 120(2):479–495
- Burer S (2012) Copositive programming. In: *Handbook on semidefinite, conic and polynomial optimization*, Springer, pp 201–218
- Burer S (2015) A gentle, geometric introduction to copositive optimization. *Math Program* 151(1):89–116
- Burer S, Dong H (2012) Representing quadratically constrained quadratic programs as generalized copositive programs. *Oper Res Lett* 40:203–206
- Calafiore GC, Ghaoui LEI (2006) On distributionally robust chance-constrained linear programs. *J Optim Theory Appl* 130(1):1–22
- Chen Z, Sim M, Xu H (2016) Distributionally robust optimization with infinitely constrained ambiguity sets. Working Paper
- Delage E, Ye Y (2010) Distributionally robust optimization under moment uncertainty with application to data-driven problems. *Oper Res* 58(3):595–612
- Dodin B (1985) Bounding the project completion time distribution in pert networks. *Oper Res* 33(4):862–881
- Duchi J, Glynn P, Namkoong H (2016) Statistics of robust optimization: a generalized empirical likelihood approach. arXiv preprint [arXiv:1610.03425](https://arxiv.org/abs/1610.03425)
- Dupačová J (1987) The minimax approach to stochastic programming and an illustrative application. *Stoch Int J Probab Stoch Process* 20(1):73–88
- Eichfelder G, Jahn J (2008) Set-semidefinite optimization. *J Convex Anal* 15:767–801
- El Ghaoui L, Oks M, Oustry F (2003) Worst-case value-at-risk and robust portfolio optimization: a conic programming approach. *Oper Res* 51(4):543–556
- Erdoğan E, Iyengar G (2006) Ambiguous chance constrained problems and robust optimization. *Math Program* 107(1–2):37–61
- Esfahani PM, Kuhn D (2015) Data-driven distributionally robust optimization using the wasserstein metric: performance guarantees and tractable reformulations. arXiv preprint [arXiv:1505.05116](https://arxiv.org/abs/1505.05116)

- Grötschel M, Lovász L, Schrijver A (1981) The ellipsoid method and its consequences in combinatorial optimization. *Combinatorica* 1(2):169–197
- Hagstrom JN (1988) Computational complexity of pert problems. *Networks* 18(2):139–147
- Halliwell LJ (2015) The lognormal random multivariate. In: Casualty actuarial society E-Forum, Spring
- Hanasusanto GA, Kuhn D (2016) Conic programming reformulations of two-stage distributionally robust linear programs over wasserstein balls. *arXiv preprint [arXiv:1609.07505](https://arxiv.org/abs/1609.07505)*
- Hanasusanto GA, Kuhn D, Wallace SW, Zymler S (2015) Distributionally robust multi-item newsvendor problems with multimodal demand distributions. *Math Program* 152(1):1–32
- Hanasusanto GA, Roitch V, Kuhn D, Wiesemann W (2017) Ambiguous joint chance constraints under mean and dispersion information. *Oper Res* 65(3):751–767
- Hu Z, Hong LJ (2013) Kullback-leibler divergence constrained distributionally robust optimization. Available at Optimization Online
- Jiang R, Guan Y (2016) Data-driven chance constrained stochastic program. *Math Program* 158(1–2):291–327
- Klabjan D, Simchi-Levi D, Song M (2013) Robust stochastic lot-sizing by means of histograms. *Prod Oper Manag* 22(3):691–710
- Lam H (2016) Recovering best statistical guarantees via the empirical divergence-based distributionally robust optimization. *arXiv preprint [arXiv:1605.09349](https://arxiv.org/abs/1605.09349)*
- Lofberg J (2004) Yalmip: a toolbox for modeling and optimization in matlab. In: 2004 IEEE international symposium on computer aided control systems design, pp 284–289
- Martello S, Toth P (1990) Knapsack problems: algorithms and computer implementations. Wiley, New York
- Möhring RH (2001) Scheduling under uncertainty: bounding the makespan distribution. In: Computational discrete mathematics, Springer, pp 79–97
- MOSEK ApS. The MOSEK optimization toolbox for MATLAB manual. Version 8.0., 2016
- Natarajan K, Teo CP (2017) On reduced semidefinite programs for second order moment bounds with applications. *Math Program* 161(1):487–518
- Natarajan K, Song M, Teo C-P (2009) Persistency model and its applications in choice modeling. *Manag Sci* 55(3):453–469
- Natarajan K, Teo CP, Zheng Z (2011) Mixed 0–1 linear programs under objective uncertainty: a completely positive representation. *Oper Res* 59(3):713–728
- Pflug G, Wozabal D (2007) Ambiguity in portfolio selection. *Quant Finance* 7(4):435–442
- Rubner Y, Tomasi C, Guibas LJ (2000) The earth mover's distance as a metric for image retrieval. *Int J Comput Vis* 40(2):99–121
- Shapiro A (2001) On duality theory of conic linear problems. In: Semi-infinite programming. Springer, pp 135–165
- Sherali HD, Adams WP (2013) A reformulation-linearization technique for solving discrete and continuous nonconvex problems, vol 31. Springer, Berlin
- Sturm JF, Zhang S (2003) On cones of nonnegative quadratic functions. *Math Oper Res* 28(2):246–267
- Vandenbergh L, Boyd S, Comanor K (2007) Generalized chebyshev bounds via semidefinite programming. *SIAM Rev* 49(1):52–64
- Wiesemann W, Kuhn D, Sim M (2014) Distributionally robust convex optimization. *Oper Res* 62(6):1358–1376
- Wolsey LA (1998) Integer programming, vol 42. Wiley, New York
- Zymler S, Kuhn D, Rustem B (2013a) Distributionally robust joint chance constraints with second-order moment information. *Math Program* 137(1–2):167–198
- Zymler S, Kuhn D, Rustem B (2013b) Worst-case value at risk of nonlinear portfolios. *Manag Sci* 59(1):172–188